

高性能计算II(B)

基于图形处理器的并行计算及CUDA编程

Ying Liu, Associate Prof., Ph.D

School of Computer and Control, University of Chinese
Academy of Sciences

Key Lab of Big Data & Knowledge Management, Chinese
Academy of Sciences

Parallel Computing


- Parallel computing platforms
- Parallel programming models
- Architecture of GPU
- Multi-threaded CPU computing vs. multi-threaded CUDA computing

Architectures

- Basic components of any architecture:
 - Processors and memory
 - Interconnect
- Logic classification based on:
 - Control mechanism
 - SISD
 - MISD
 - SIMD (Single Instruction Multiple Data stream)
 - MIMD (Multiple Instruction Multiple Data stream)
 - Address space organization
 - Shared Address Space
 - Distributed Address Space

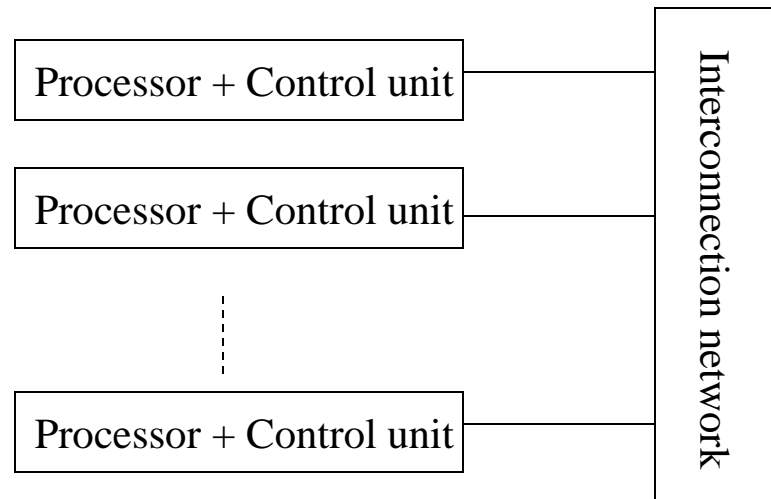
SIMD Architecture

- Multiple processors execute the same instruction
- Data that each processor sees may be different
- Individual processors can be turned on/off at each cycle (“masking”)
- Examples: Illiac IV, Thinking Machines’ CM-2, DAP, ...
- Specialized architecture limits the popularity




MIMD Architecture

- Each processor executes program independent of other processors
- Processors operate on separate data streams
- May have separate clocks
- Examples: IBM SP, Cray T3D & T3E, SGI Origin, Clusters, Sun Ultra Servers, ...




Shared Address Space

- All processors share a single global address space
- Single address space facilitates a simple programming model
- Examples: SGI Origin 2000, IBM SP2






(a) UMA



(b) NUMA


Static Interconnects

- Consist of point-to-point links between processors
- Examples: hypercube, mesh/torus, fat tree
- Two major characteristics:
 - Can make parallel system expansion easy
 - Some processors may be “closer” than others




Multidimensional Meshes

- Each processor in a d -dimensional mesh is connected to $2d$ other processors
- Data between processors routed via intermediate processors
- In practice, only 2 or 3 dimensional meshes are constructed
- Mesh with wrap around: Torus




Hypercube

- A hypercube is a multidimensional mesh with exactly 2 processors in each dimension
- In a d -dimensional hypercube, each processor is connected with d other processors




Fat Tree

- A tree network is one where there is exactly one path between a pair of processors
- In a simple tree, the bandwidth on higher level links is shared among all processors below creating bottlenecks
- A fat tree solves the problem by using wider links at higher levels in the tree



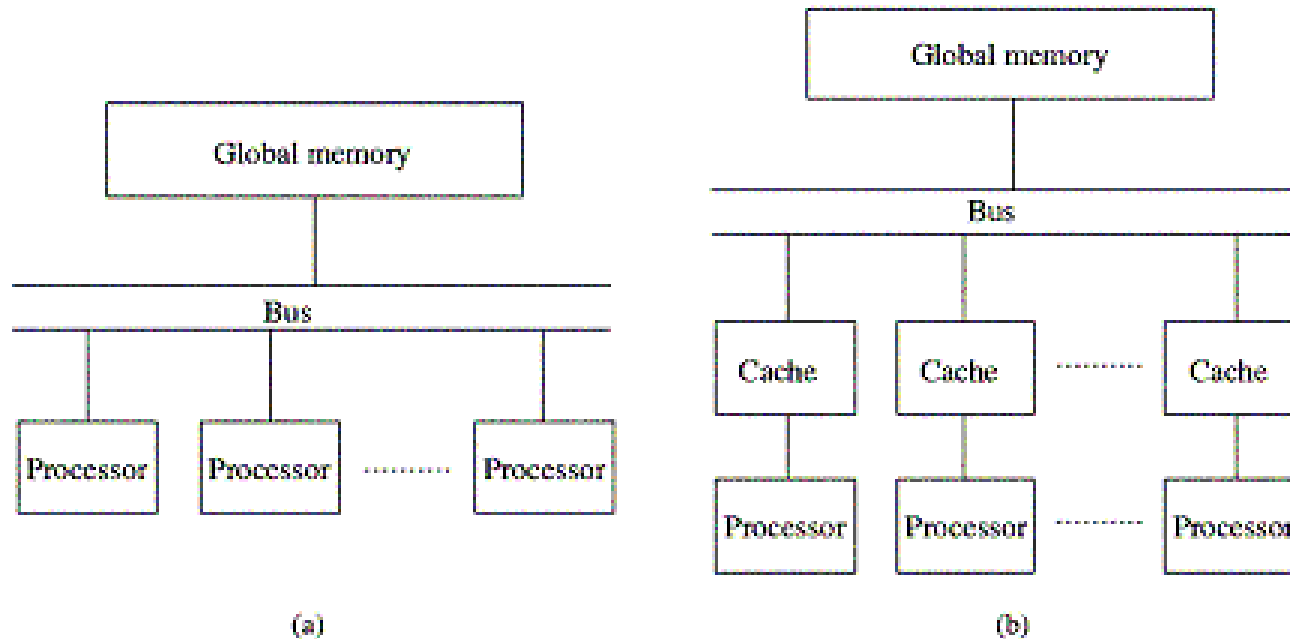
Dynamic Interconnects

- Paths are established as needed between processors
- Examples: bus based, crossbar, multistage networks
- Two major characteristics:
 - System expansion difficult
 - Processors are usually equidistant




Bus Systems

- Only one pair of processors or processor and memory can exchange data at a time
- Shared medium, good for broadcasting
- Bandwidth limitation, limited to dozens of nodes



Crossbars

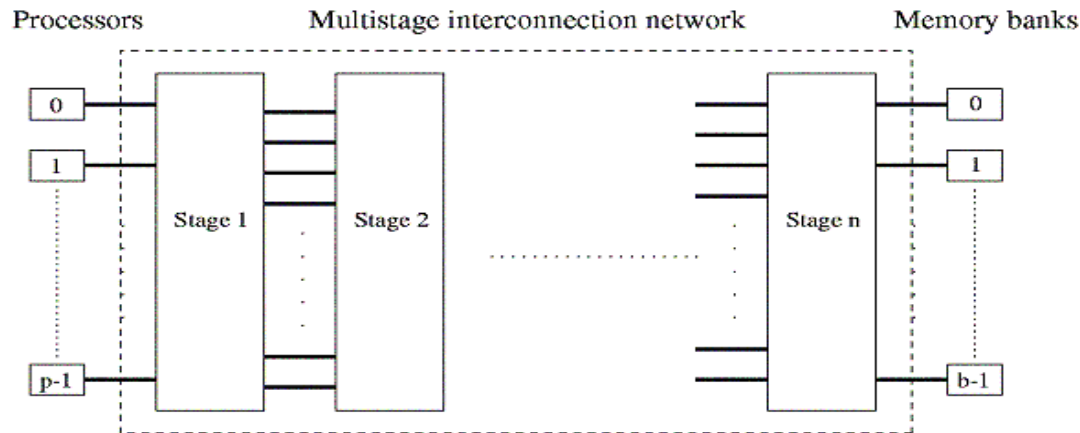
- Designed to increase the traffic between processors and memory
- Given N processors and an N -way interleaved memory, an $N \times N$ crossbar provides N simultaneous connections
- Can control crossbar settings to achieve any desired permutation



An 8×8 crossbar switch ($8! = 40320$ permutations)

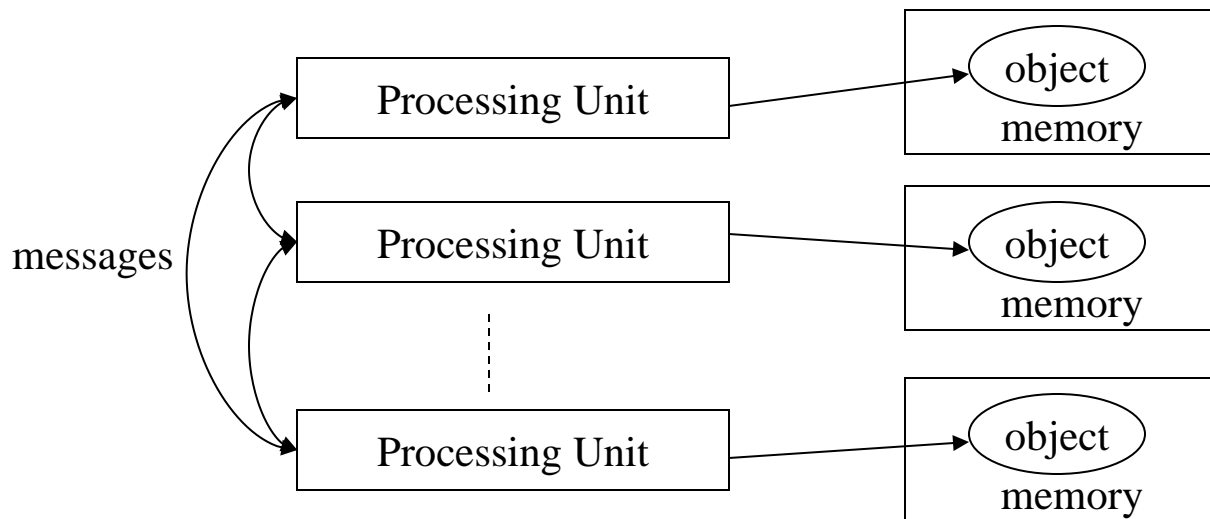
Multistage Network Systems

- Multistage networks use $N \log(N)$ switches
- Use multiple stages of switches to build interconnects
- Design choices in multistage networks:
 - Basic switch type and size
 - Number of stages
 - Connections between stages



Message Passing Platform (Distributed Address Space)

- “Shared nothing:” each processor has a private memory
- Processors can directly access only local data
- Each processing unit can be single processor or a multiprocessor
- Interaction between processors relies on message passing
- Example: clusters



Operating Systems

- Need to support tasks similar to serial OS like UNIX
 - Memory and process management
 - File system
 - Security
- Additional support needed
 - Job scheduling: assign tasks to idle processors, time sharing, space sharing
 - Parallel programming support: message passing, synchronization

Parallel Computing

- Parallel computing platforms
- **Parallel programming models**
- Architecture of GPU
- Multi-threaded CPU computing vs. multi-threaded CUDA computing

What is Parallel Programming?

- Parallel programming involves constructing or modifying a program for solving a given problem on a parallel machine
 - Start from a serial algorithm for solving the problem
- Goals of parallel programming – Performance!
- Effective parallel programming
 - *Key1: Minimization of inter-processor synchronization costs*
 - *Key 2: Equal workload between processors*

Alternatives to Parallel Programming

- Automatic parallelizing compiler
- Libraries
- Implicit parallel programming:
 - Programmer directives to help compiler analysis
 - Example OpenMP
- Explicit parallel programming
 - Easiest to support
 - Example MPI

Types of Parallelism

■ Data parallelism:

- Partition the data across processors
- Each processor performs the same operations on its local data partitioning

■ Task parallelism:

- Assign independent modules to different processors
- Each processor performs different operations

Data Dependence & Parallelization

- Consider the following loop of a C program

```
for (i=0;i<1000;i++)  
    a[i]=b[i]+c[i]
```

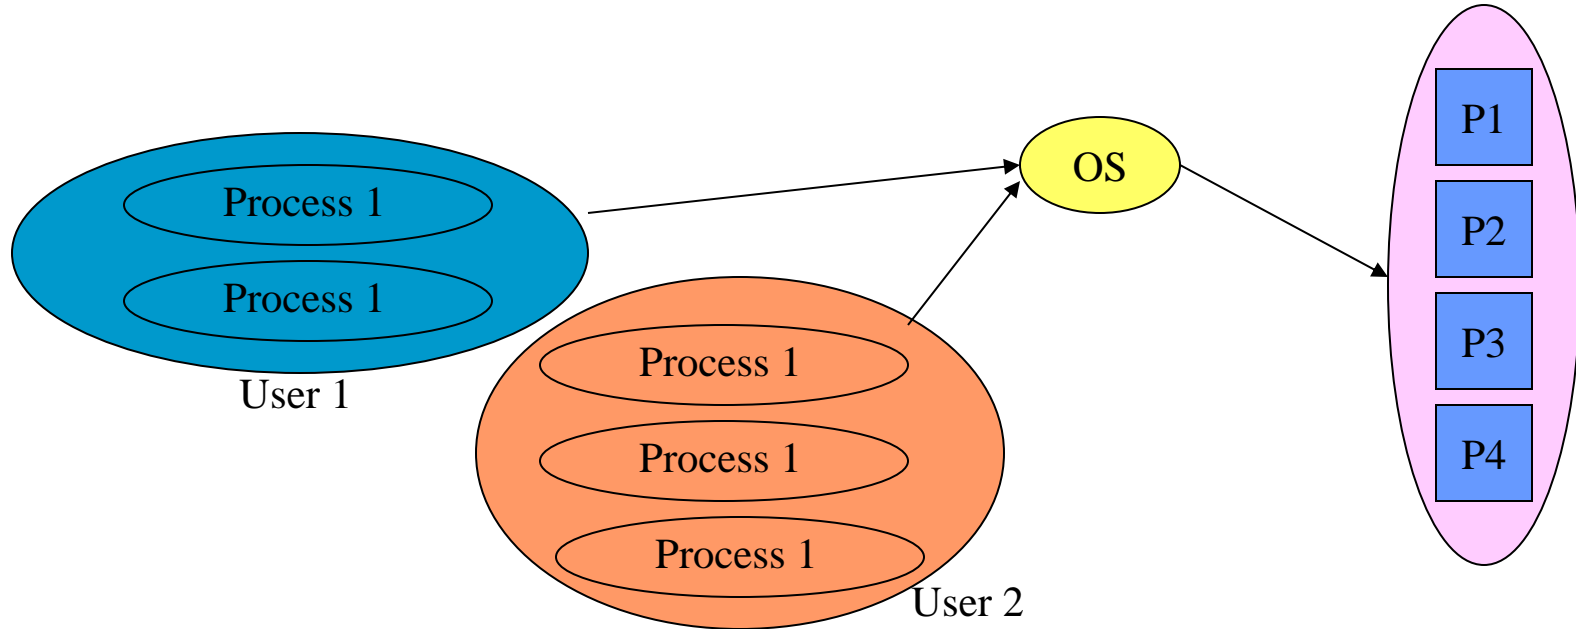
- If one unfolds the loops, the statements would be executed as follows:

```
a[0]=b[0]+c[0];  
a[1]=b[1]+c[1];  
.....  
a[999]=b[999]+c[999];
```

- Can each iteration be executed in parallel?

Processes


- A process is a program along with all of its environment
 - How to communicate?
 - Who creates it?
 - Who map processes to physical processors?



Process Creation

- Use a process creation primitive:

```
m_set_procs(nproc);  
m_fork(func,[args]);
```




- Each child process is an exact copy of its parent
- All the variables associated with the parent are private for the parent and each child unless explicitly made shared
- The parent had ID=0, children share the other $nproc-1$ IDs

```
id=m_get_myid();  
nprocs=m_get_procs();
```

Process Destruction

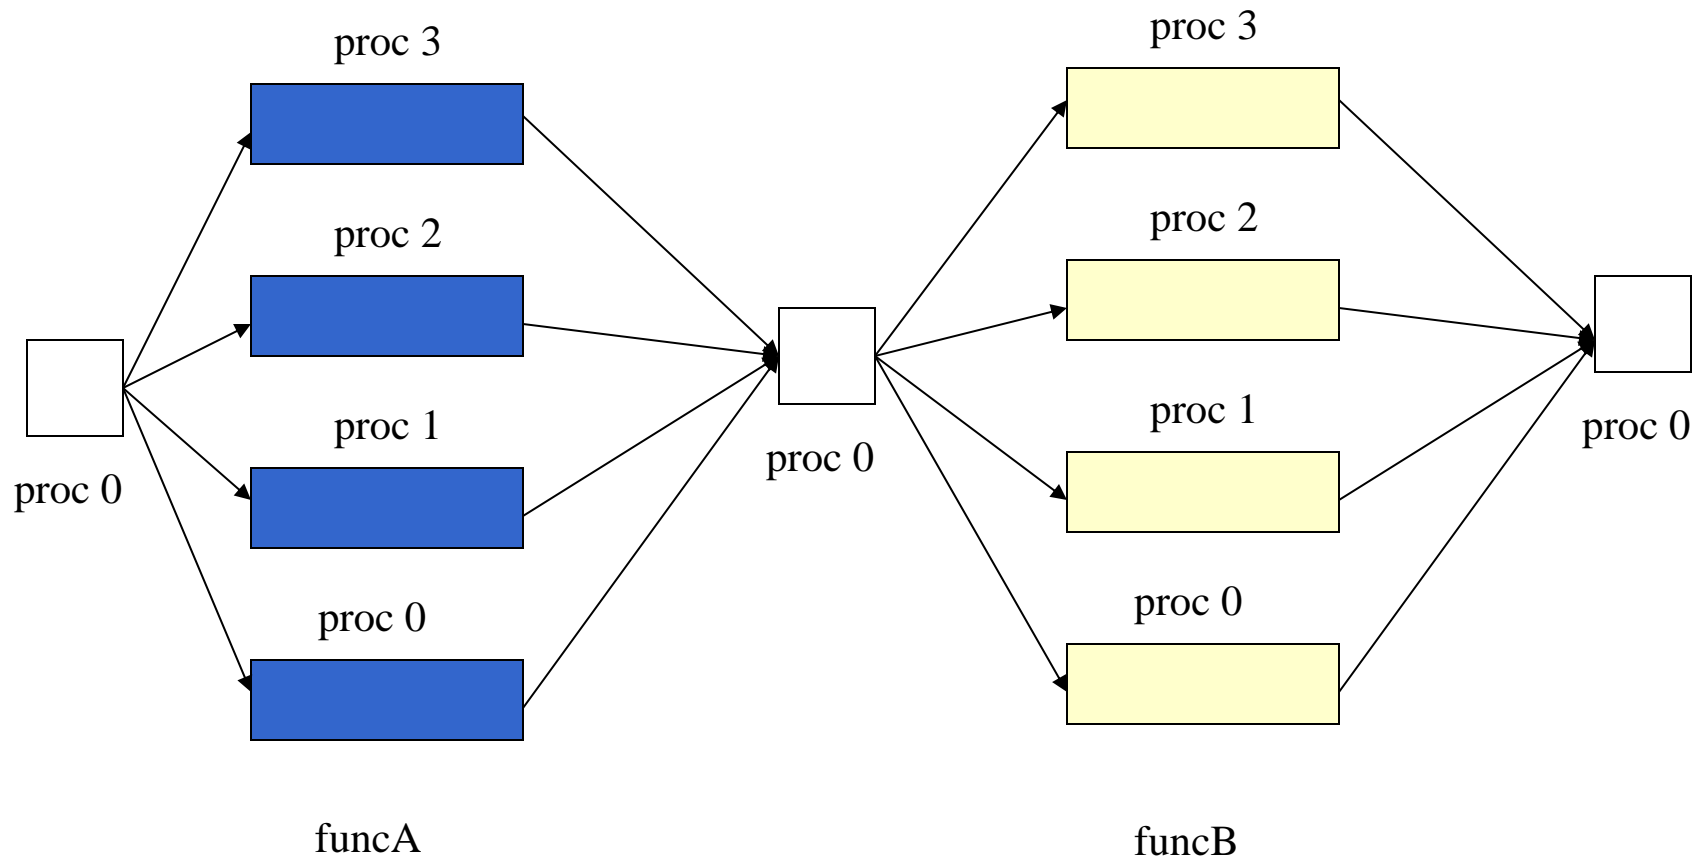
- All child processes are terminated; only the parent remains
- The parent waits for all child processes to terminated before returning
- Use a process destruction primitive

`m_kill_procs();`




Parallel Program Model


- Processes executing in a data parallel fashion




Loop Scheduling




Prescheduling



Static blocked



Static interleaved



Dynamic

Prescheduling

- Assign work using a schedule array

```
float *a, *b, *c;
```

```
int low[ ] = {0,50,80}
```

```
int high[ ] = {50,80,120}
```

```
void parallel_func( ) {
```

```
    int id, i;
```

```
    id = m_get_myid( );
```

```
    for (i=low[id]; i<high[id]; i++)
```

```
        a[i] = b[i] + c[i];
```

```
}
```

Static Blocked Scheduling

- Assign a contiguous chunk of iterations based on process ID

```
float *a, *b, *c;
void parallel_func( ) {
    int id, i, nprocs;
    int low, high;
    id = m_get_myid( );
    nprocs = m_get_numprocs( );
    low = id*100 / nprocs;
    high = (id+1)*100 / nprocs;
    for (i=low; i<high; i++)
        a[i] = b[i] + c[i];
} /* end of parallel_func */
```

Static Interleaved Scheduling

- Assign iterations in a round-robin way based on process ID

```
float *a, *b, *c;  
void parallel_func( ) {  
    int id, i, nprocs;  
    id = m_get_myid( );  
    nprocs = m_get_numprocs( );  
    for (i=id; i<100; i+=nprocs)  
        a[i] = b[i] + c[i];  
}
```

Dynamic Scheduling

- Self scheduling: processes execute iterations using a shared counter

```
float *a, *b, *c;
int i;

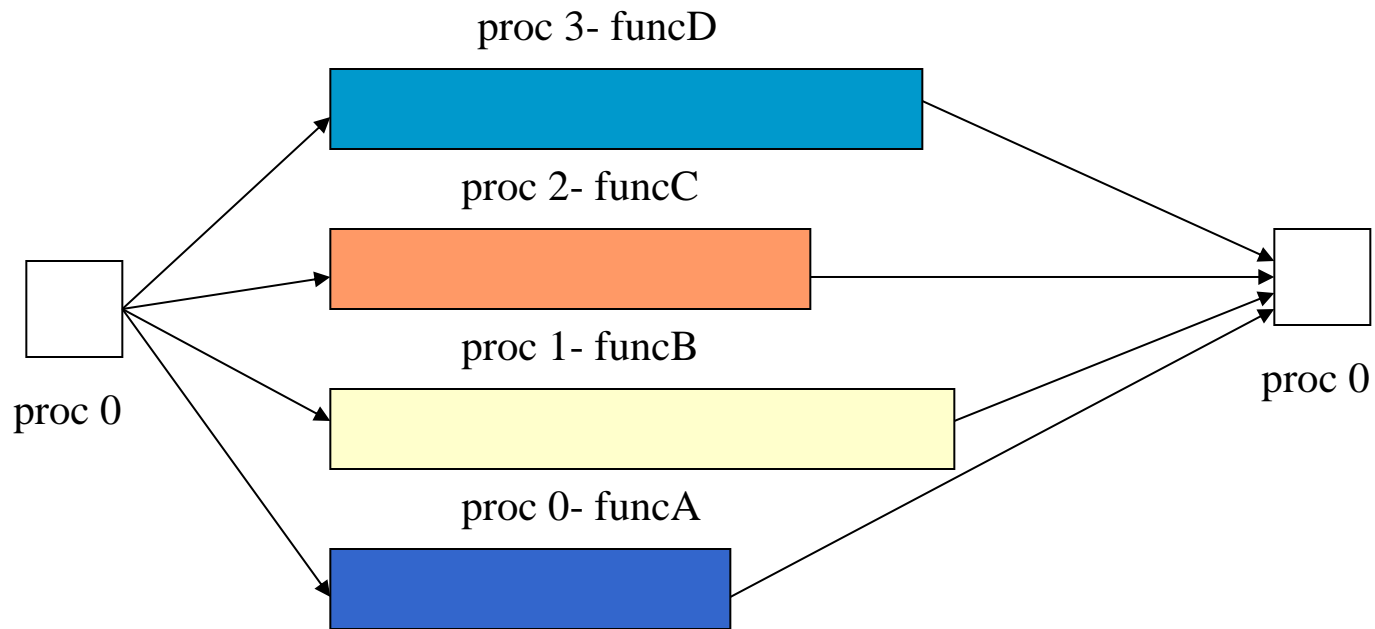
i = 0;

void parallel_func() {
    int i1, i2;
    int chunk=2;

    while (i1<100) {
        m_lock( );
        i1 = i;
        i += chunk;
        m_unlock( );
        for (i2=i1; i2<min(i1+chunk,100); i2++)
            a[i2] = b[i2] + c[i2];
    } /* end of while */
} /* end of parallel_func */
```

Parallel Program Model

- Processes executing in a function/task parallel fashion

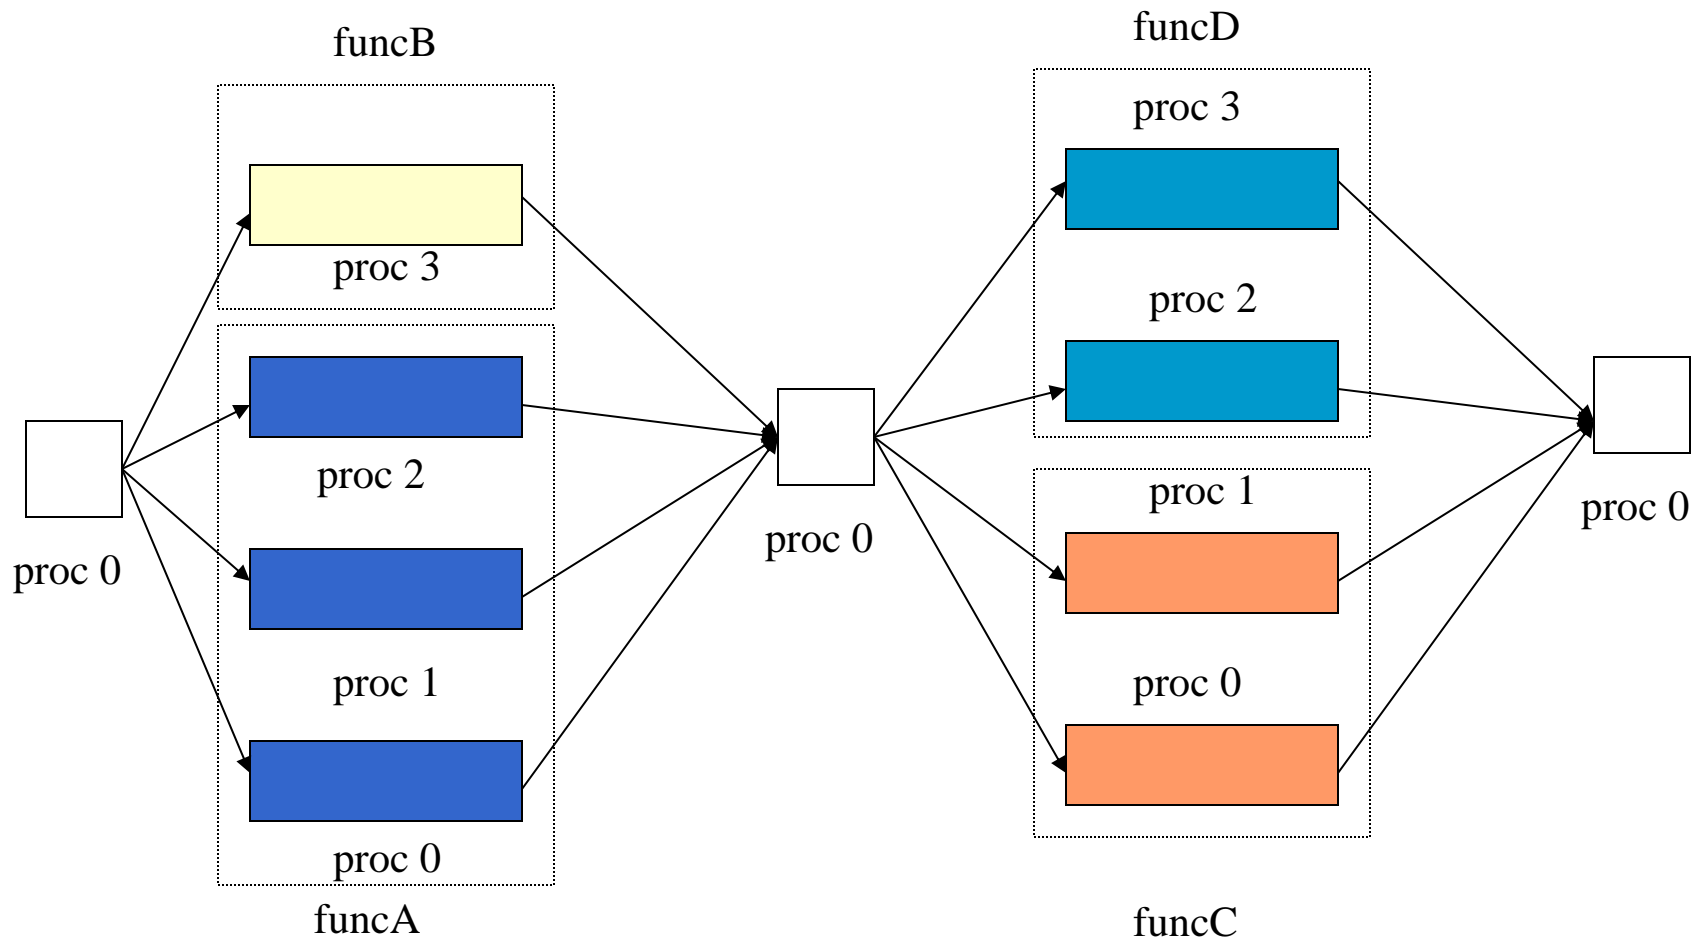


Parallel Program Model

```
main() {  
    m_set_procs(4);  
    m_fork(func);  
    m_kill_procs();  
}  
void func() {  
    Int id;  
    id = m_get_myid();  
    switch (id) {  
        case 0: funcA();  
            break;  
        case 1: funcB();  
            break;  
        case 2: funcC();  
            break;  
        case 3: funcD();  
            break;  
    }  
}
```


Parallel Program Model

- Processes executing in a function & data parallel fashion



Parallel Program Model

```
main() {  
    m_set_procs(4);  
    m_fork(func1);  
    m_park_procs();  
    .....  
    m_rele_procs();  
    func2();  
    m_kill_procs();  
}
```

```
void func1() {  
    switch (id) {  
        case 0,1,2: funcA();  
                    break;  
        case 3: funcB();  
                break;  
    }  
}  
  
void func2() {  
    switch (id) {  
        case 0,1: funcC();  
                break;  
        case 2,3: funcD();  
                break;  
    }  
}
```

Performance of Parallel Programs

- T = execution time of the best serial algorithm
- T_p = execution time of the parallel algorithm with p processors
- Speedup, $S_p = T / T_p$
- Efficiency $E_p = S_p / p$
- However, E_p is realistically never 100% due to:
 - Synchronization and communication cost across processors
 - Less-than-perfect load balancing among processors
 - Assume also that α represents the fraction that cannot be parallelized, i.e., the serial fraction
 - Then, using p processors in parallel and assuming perfect speedups in the parallelized portion, the parallel execution time,

$$T_p = T (\alpha + (1 - \alpha) / p)$$

Example

```
float sum,sum0,sum1;
main() {
    m_set_procs(2);
    m_fork(parallel_func);
    m_kill_procs();
    sum=sum0+sum1;
    printf("total sum is %lf\n",sum);
}
void parallel_func(){
    int id;
    id=m_get_myid();

    if (id==0) sum0=1.0+2.0;
    else sum1=3.0+4.0;
}
```

Example

```
float total_sum;
main() {
    total_sum=0.0;
    m_set_procs(2);
    m_fork(parallel_func);
    m_kill_procs();
    printf("total sum is %f\n",total_sum);
}
void parallel_func() {
    int id;
    float partial_sum;
    id=m_get_myid();

    if (id==0) partial_sum=1.0+2.0;
    else partial_sum=3.0+4.0;

    total_sum+=partial_sum;
```

```
}
```

Example

Possible outputs?

Total sum is 10.0

(ts = 0;

P0: ts = ts(0) + 3;

P1: ts = ts(3) + 7;

Total sum is 7.0

(ts = 0;

P0: ts = ts(0) + 3;

P1: ts = ts(0) + 7;

Total sum is 3.0

(ts = 0;

P1: ts = ts(0) + 7;

P0: ts = ts(0) + 3;

Contention

- Random process scheduling creates problems with respect to modifying shared variables
- The final values of a shared variable can differ from run to run depending on the order in which it was modified
- Root of this contention problem is the simultaneous accessing of a shared variable
- Solution ?

Restrict the access to shared variables

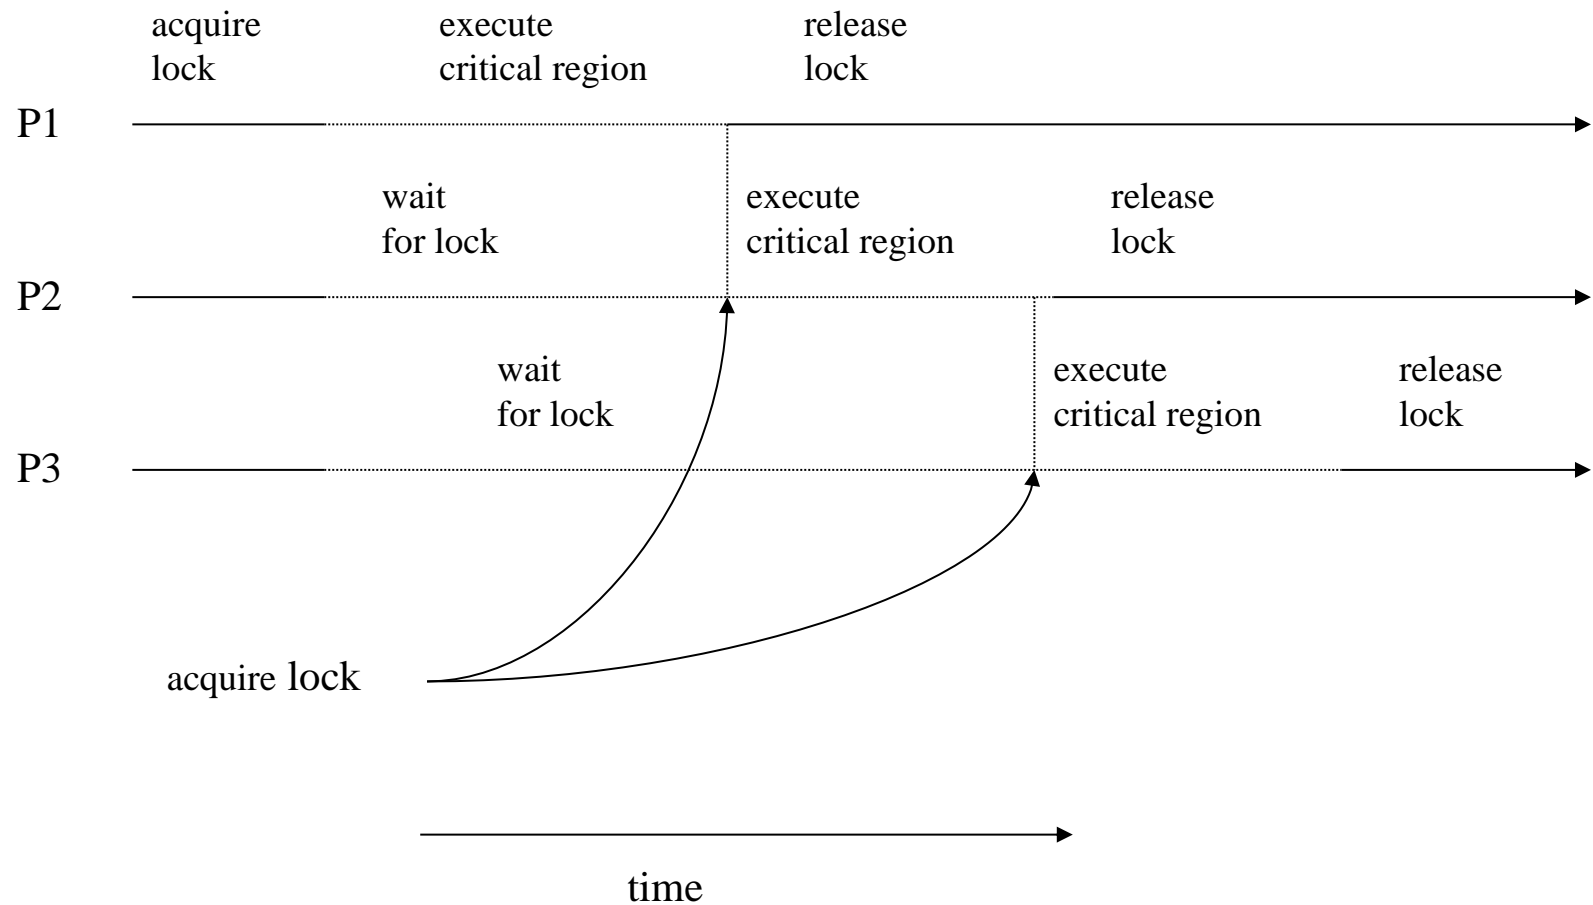
Variable Locks

- Locks are used to provide exclusive access for the modification of a variable
- A lock variable is to be created for every shared variable that needs exclusive access:
 - It must be shared among all processes that need to use it
 - It can either be in a “locked” or “unlocked” state

Variable Locks

- The steps to be followed for exclusive access using locks are:
 - Acquire the lock for a variable
 - Modify the variable
 - Release the lock for the variable
- The use of locks sequentializes execution of program sections and hurts performance

Illustration of Locks



Example

```
float total_sum;
void parallel_func() {
    int id;
    float partial_sum;
    id=m_get_myid();

    if (id==0) partial_sum=1.0+2.0;
    else partial_sum=3.0+4.0;

    m_lock();
    total_sum+=partial_sum;
    m_unlock();
}
```

Race Conditions

```
float total_sum;
void parallel_func() {
    int id;
    float partial_sum;
    id=m_get_myid();

    if (id==0) partial_sum=1.0+2.0;
    else partial_sum=3.0+4.0;

    m_lock();
    total_sum+=partial_sum;
    m_unlock();

    average=total_sum/4.0;
    printf("%d average is %f\n",id, average);
}
```

Race Conditions

Possible output

0 average is 2.5

1 average is 2.5

1 average is 1.75

0 average is 2.5

0 average is 0.75

1 average is 2.5

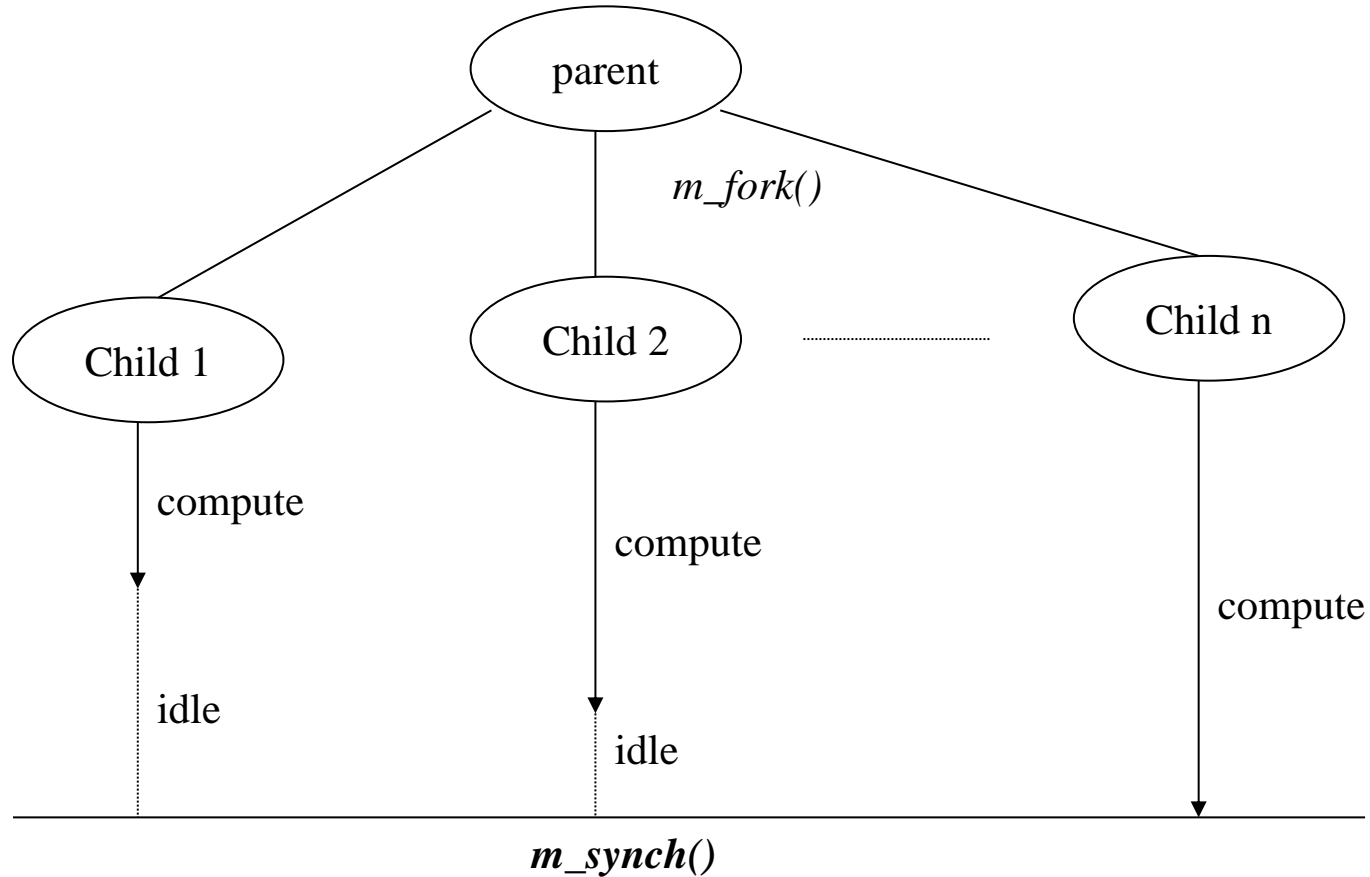
Barrier

- A race condition exists when?

The results of a parallel program depend on the relative execution speed of processes

- Barriers enable processes to synchronize with each other and help avoid race conditions
- When a process enters a barrier, it waits for all other processes involved to reach the barrier before continuing
- Barriers cause performance degradation and therefore must be used carefully

Illustration of Barrier



Barriers

```
float total_sum;
void parallel_func() {
    int id;
    float partial_sum;
    id=m_get_myid();
    if (id==0) partial_sum=1.0+2.0;
    else partial_sum=3.0+4.0;
    m_lock();
    total_sum+=partial_sum;
    m_unlock();

    m_synch( );

    average=total_sum/4.0;
    printf("%d average is %f\n",id, average);
}
```


Process Summary

■ Contention

- Locks for exclusive access to shared variables

■ Race conditions

- Barriers for synchronization

Loop Parallelization

- We will look at some simple and commonly occurring code fragments and consider how to parallelize them
- Techniques involve to remove dependencies
 - Local variables to remove dependencies
 - Code transformations
 - Use of particular scheduling techniques

Local Variables

■ Code Fragment

```
float *a, *b, x;  
int n;
```

```
for (i=0 ;i<n; i++) {  
    x = a[i]*3;  
    b[i] = x*b[i];  
}
```

Local Variables

```
float *a, *b;
```

```
int n;
```

```
void parallel_func(n) {
```

```
    int i, id, nprocs;
```

```
    float x;
```

```
    id = me_get_myid();
```

```
    nprocs = m_get_numprocs();
```

```
    for (i=id ;i<n;i+=nprocs) {
```

```
        x=a[i]*3;
```

```
        b[i]=x*b[i];
```

```
    }
```

```
}
```

Loop-Carried Variables

- Code fragment

```
indx=0;  
for (i=0;i<n;i++) {  
    indx+=i;  
    a[i]=b[indx];  
}
```

Loop-Carried Variables

```
float *a, *b;
int n;

void parallel_func(n) {
    int i, id, nprocs, indx;

    id = me_get_myid();
    nprocs = m_get_numprocs();
    for (i=id ;i<n;i+=nprocs) {
        indx = (i*(i+1))/2;
        a[i]=b[indx];
    }
}
```

Indirect Indexing

- Code fragment

```
for (i=0;i<n;i++) {  
    ix=indexx[i];  
    iix=ixoffset[ix];  
    total[iix]=total[iix]+delta;  
}
```

Indirect Indexing

```
for (i=0;i<n;i++) {  
    ix=indexx[i];  
    iix [i]=ixoffset[ix];  
}  
  
for (i=0;i<n;i++) {  
    total[iix[i]]=total[iix[i]]+delta;  
}
```


Loop Reordering

- Code fragment

```
for (k=0;k<n;k++) {  
    for (i=0;i<n;i++) {  
        for (j=0;j<n;j++) {  
            a[i][j]+=b[i][k]+c[k][j];  
        }  
    }  
}
```

Loop Reordering

```
float *a,*b,*c;
void parallel_func() {
    int i, id, nprocs;

    id = m_get_myid();
    nprocs = m_get_numprocs();
    for (i=id; i<n; i+=nprocs)
        for (k=0; k<n; k++)
            for(j=0; j<n; j++)
                a[i][j]+=b[i][k]+c[k][j];
}
```

Processes vs. Threads

■ Process

- a process may have multiple threads
- heavyweight
- allocated a private memory region

■ Thread

- share memory region with its parent
- lightweight
- share CPU time slots

Parallel Program Model Summary

- Data parallel programming dominates
- Loops in programs are the source of data parallelism
- Exploitation of parallelism involves sharing work in loops among processes
- Have to use appropriate scheduling techniques for optimal work sharing
- Parallelism in loops not always straightforward to find due to dependence
- Have to perform some transformations to expose parallelism

OpenMP

- A standard for directive based parallel programming
- Higher level than Pthreads, no threads manipulating
- API used in FORTRAN, C, C++
- Support for concurrency, synchronization, etc.

OpenMP Programming Model

- #pragma omp directive [clause list]

- #pragma omp parallel [clause list]

/*structured block*/

- Example

```
int a, b;
```

```
main() {
```

```
    //serial segment
```

```
    #pragma omp parallel num_threads (8) private (a) shared (b)
```

```
    {
```

```
        //parallel segment
```

```
    }
```

```
    //serial segment
```

```
}
```

Message Passing Interface (MPI)

- MPI – standard for explicit message passing (C and Fortran)
- Portable
- Widely used, requires minimal underlying hardware knowledge
- Developed by a group of researchers
- Support almost all hardware vendors


MPI Primitives

- MPI has over 125 functions, but need to know only 6 to get started on writing real applications
 - MPI_Init: initiate an MPI computation
 - MPI_Finalize: terminate an MPI computation
 - MPI_Comm_Size: determine number of processes
 - MPI_Comm_Rank: determine my process identifier
 - MPI_Send: send a message
 - MPI_Recv: receive a message


Parallel Computing

- Parallel computing platforms
- Parallel programming models
- **Architecture of GPU**
- Multi-threaded CPU computing vs. multi-threaded CUDA computing

GPU Design Philosophy




CPU



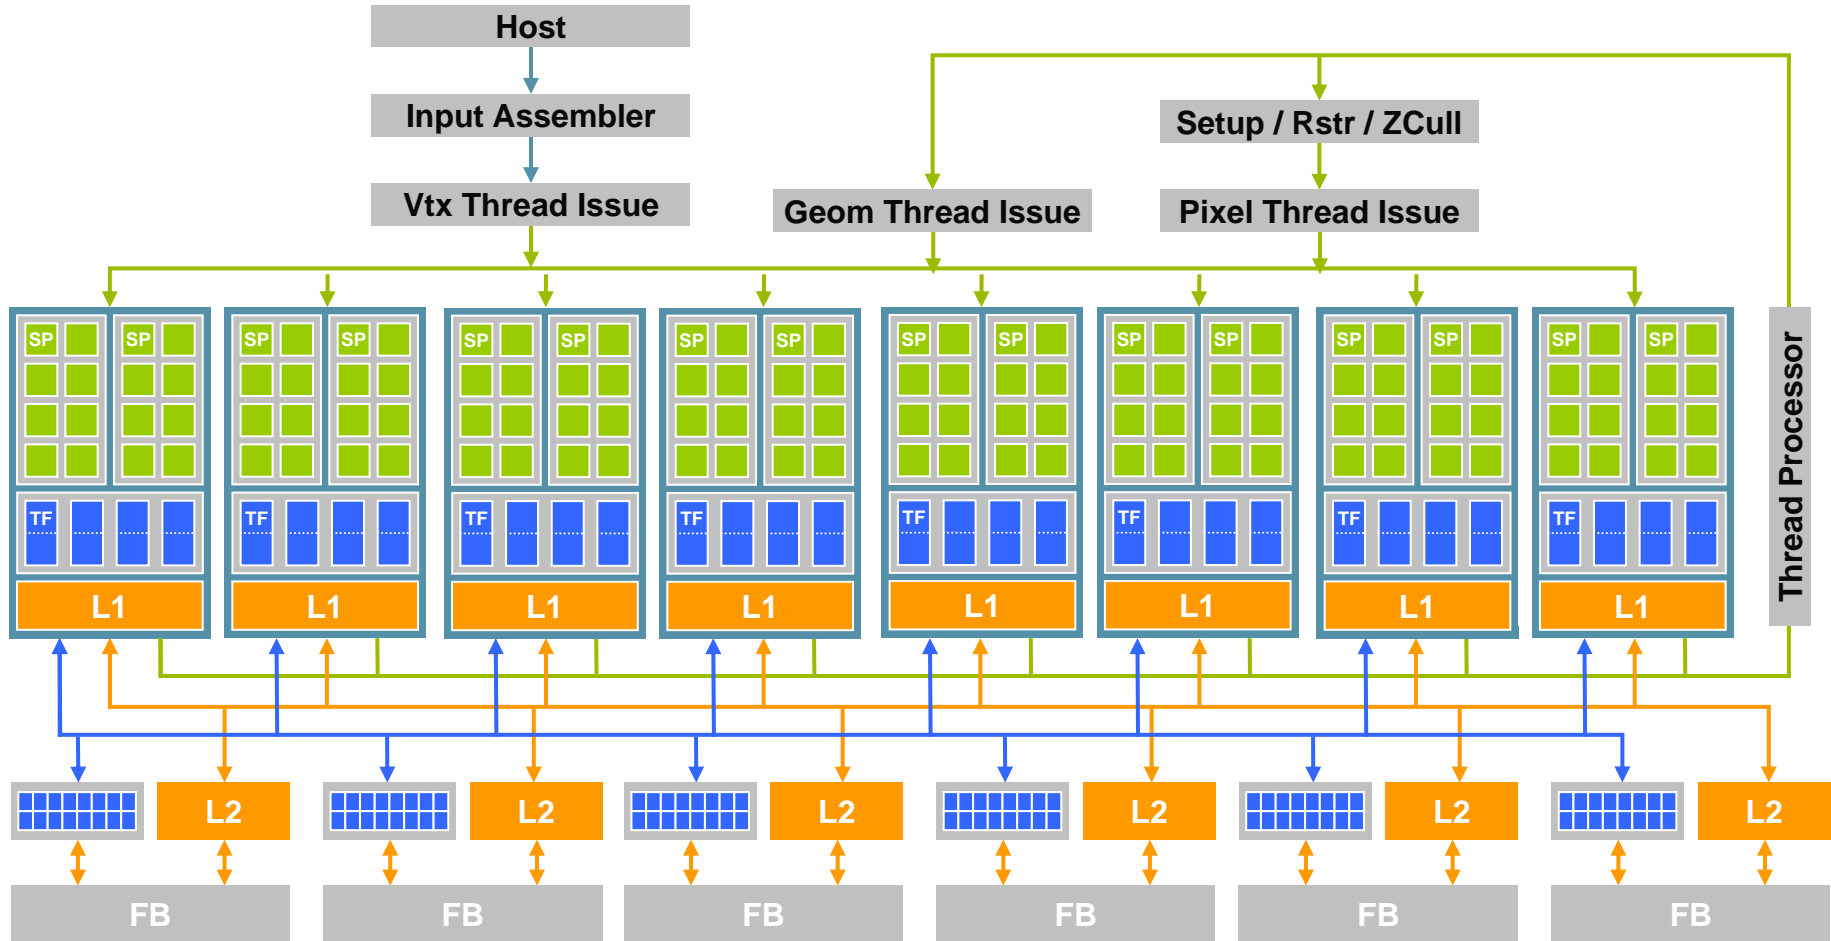
GPU

- Difference between GPU and CPU
 - More transistors for data processing
 - Many-core (hundreds of cores)

Nvidia's GeForce 8800

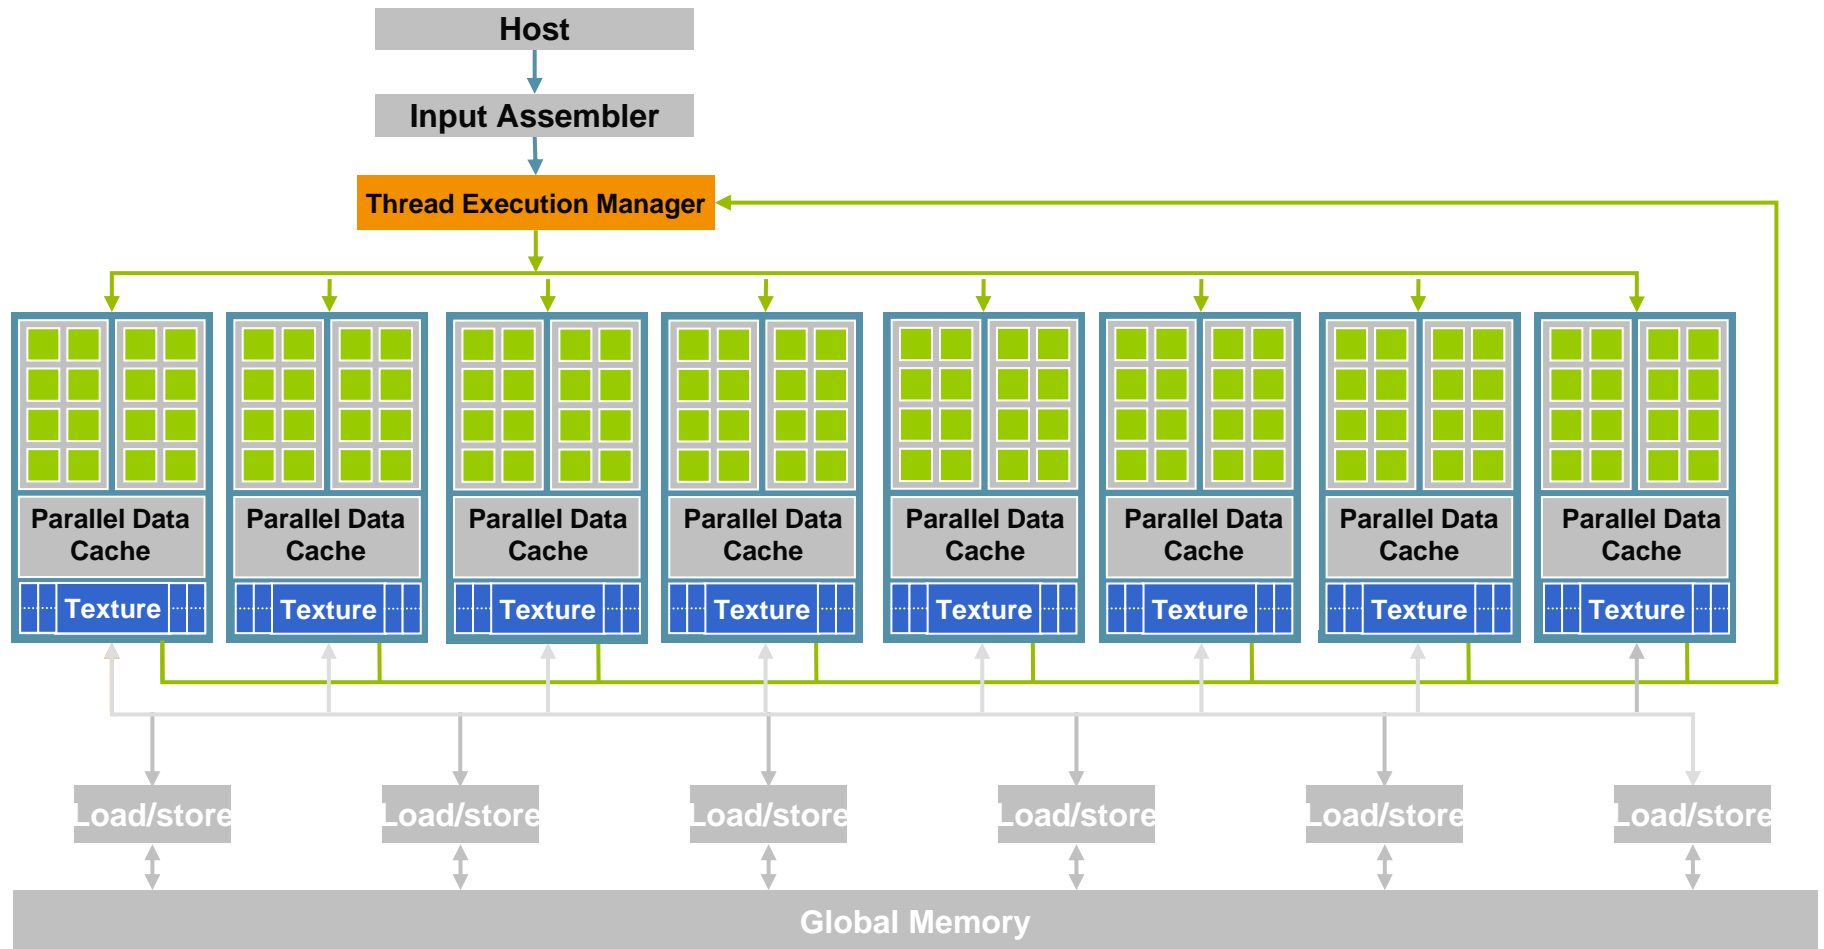


G80 – Graphics Mode



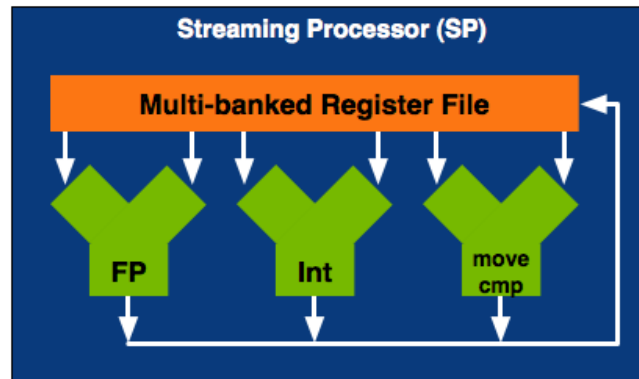
Block Diagram of the GeForce 8800

G80 CUDA Mode



Streaming Processor (SP)


- A fully pipelined, single-issue, in-order microprocessor
 - 2 ALUs and a FPU
 - Register file
 - 32-bit scalar processing
 - No cache



Streaming Multiprocessor (SM)

■ An array of SPs

- 8 streaming processors
- 2 Special Function Units (SFU)
 - Transcendental operations (e.g. sin, cos) and interpolation
- A 16KB read/write shared memory
 - Not a cache, but a software-managed data store
- Multithreading issuing unit
 - Dispatch instructions
- Instruction cache
- Constant cache



Parallel Computing

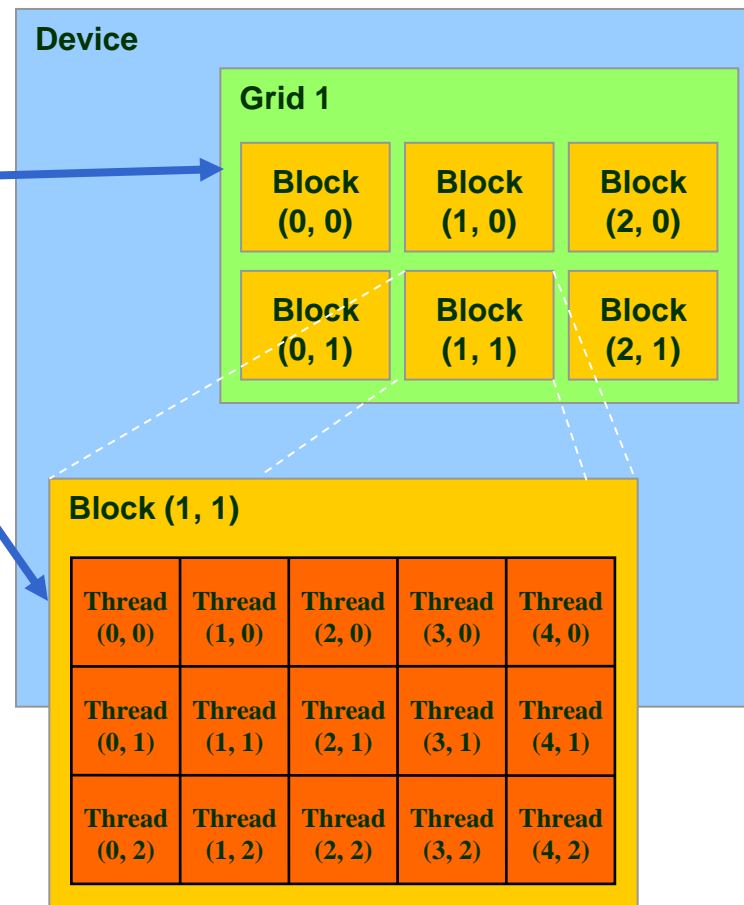
- Parallel computing platforms
- Parallel programming models
- Architecture of GPU
- Multi-threaded CPU computing vs. multi-threaded CUDA computing

CUDA Device

- A compute **device**
 - Is a co-processor to the CPU or **host**
 - Has its own DRAM (**device memory**)
 - Runs many **threads in parallel**
 - Is typically a **GPU** but can also be another type of parallel processing device
- **Kernel** — Data-parallel portions of an application which run on many threads

Block IDs and Thread IDs

- Each thread uses IDs to decide what data to work on
 - Block ID: 1D or 2D
 - Thread ID: 1D, 2D, or 3D
- Simplify memory addressing when processing multidimensional data
 - Image processing
 - Solving PDEs on volumes
 - ...



CUDA


■ Integrated host+device C program

- Serial or modestly parallel parts in host C code
- Highly parallel parts in device SPMD kernel C code

Serial Code (host)

Parallel Kernel (device)


`KernelA<<< nBlk, nTid >>>(args);`



Serial Code (host)

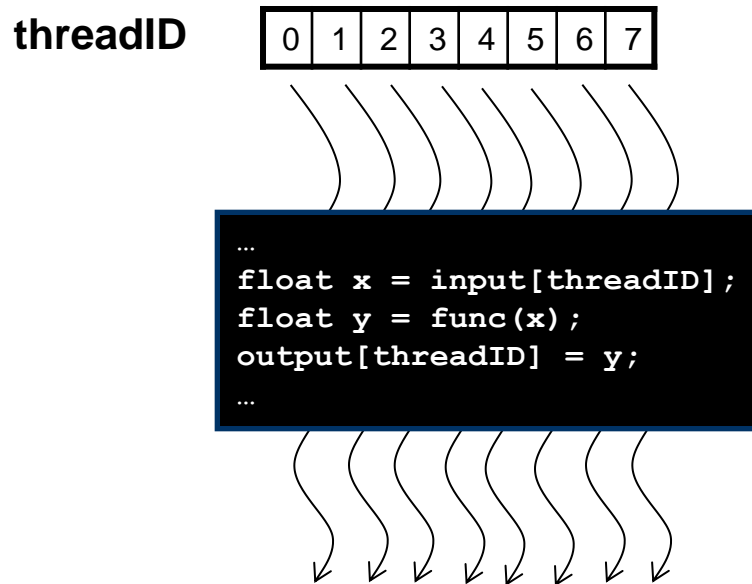
Parallel Kernel (device)

`KernelB<<< nBlk, nTid >>>(args);`



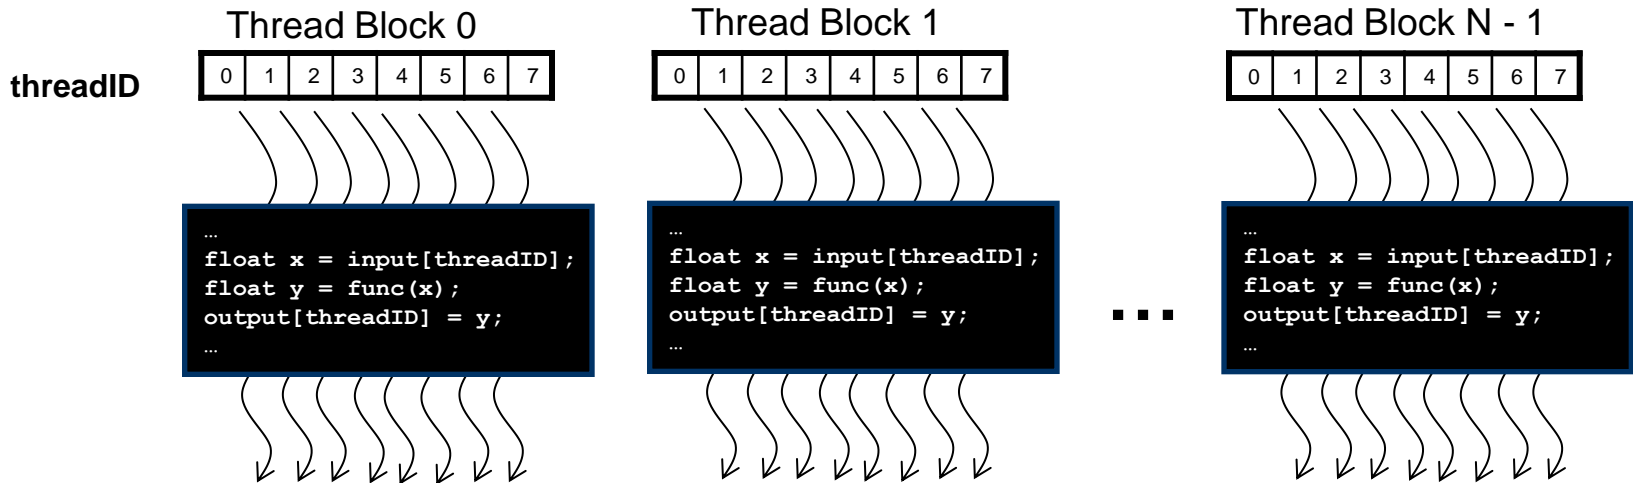
Arrays of Parallel Threads

- A CUDA kernel is executed by an array of threads
 - All threads run the same code (SPMD)
 - Each thread has an ID that it uses to compute memory addresses and make control decisions



Thread Blocks

- Divide thread array into multiple blocks
 - Threads within a block cooperate via *shared memory*, *atomic operations* and *barrier synchronization*
 - Threads in different blocks cannot cooperate



Differences Between GPU Threads and CPU Processes

- GPU threads execute on streaming processors (SPs)
 - CPU threads (processes) execute on processors
- GPU threads are extremely lightweight
 - Very little creation overhead
- GPU needs 1000s of threads for full efficiency
 - Multi-core CPU needs only a few
- GPU good at applications of little communication between threads

GPU + CPU Cluster

■ Host

- C/C++
- MPI primitives
- One CPU core is explicitly coupled with a GPU
- GPU kernel is launched by a CPU process

CUDA API Highlights: Easy and Lightweight

- The API is an extension to the ANSI C programming language
 - Low learning curve
- The hardware is designed to enable lightweight runtime and driver
 - High performance